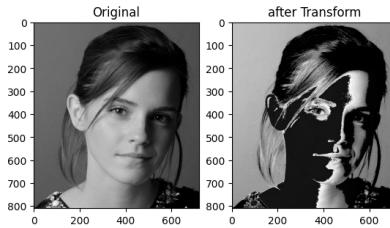


EN3160: Intensity Transformations and Neighborhood Filtering

1. Implement the intensity transformation depicted in Fig. 1 on the image

```

1 # for section 1 - linear transformation with unit gradient
2 section1 = img < 50
3 output[section1] = img[section1]
4 #section 2
5 section2= (img >= 50) & (img < 150)
6 output[section2]= (15.5 * img[section2]) + 50
7 # for section three - linear transformation with unit gradient
8 section3 = (img >= 150) & (img < 255)
9 output[section3]= img[section3]
```



2. Apply a similar operation as above (question 1) to accentuate (a) white matter and (b) gray matter

White Matter (WM) appears brighter in T1-weighted MRI (high intensity values). Gray Matter (GM) appears darker than WM but brighter than CSF (medium intensity values). We can threshold or remap these intensity ranges to make them stand out.

```

1 # Normalize image for consistent scaling to the range 0-255
2 img_norm = cv2.normalize(img2, None, 0, 255, cv2.NORM_MINMAX)
3 # White Matter mask (bright intensities, e.g., 180 255)
4 wm_mask = cv2.inRange(img_norm, 180, 255)
5 # Gray Matter mask (medium intensities, e.g., 100 179)
6 gm_mask = cv2.inRange(img_norm, 100, 179)
7
8 # Apply masks (accentuate by making them brighter)
9 wm_highlight = cv2.bitwise_and(img_norm, img_norm, mask=wm_mask)
10 gm_highlight = cv2.bitwise_and(img_norm, img_norm, mask=gm_mask)
11
12 # Optionally enhance contrast
13 wm_highlight = cv2.convertScaleAbs(wm_highlight, alpha=1.5, beta=0)
14 gm_highlight = cv2.convertScaleAbs(gm_highlight, alpha=1.5, beta=0)
```

The mask determines which output pixels are calculated and which are set to zero. Only pixels where mask = 255 are taken from (img and img). All other positions become black (intensity 0).

3. Apply gamma correction to the L plane in the Lab color space and state the γ value.

What is the L*a*b Color Space? The L*a*b (CIELAB) color space is designed to be **perceptually uniform** meaning changes in values correspond more closely to how humans perceive color differences. It separates **lightness** from

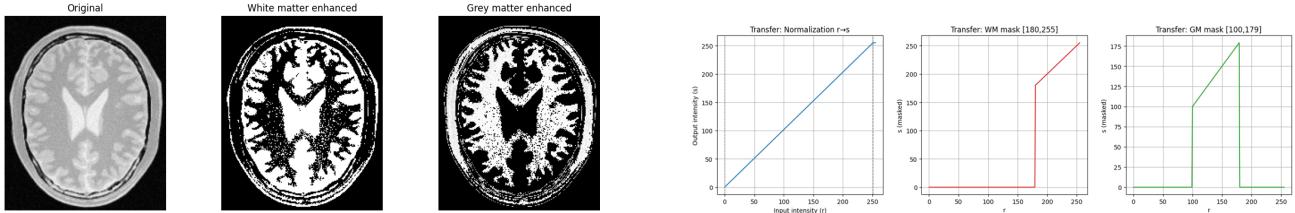


Figure 1: Enter overall caption here describing both images.

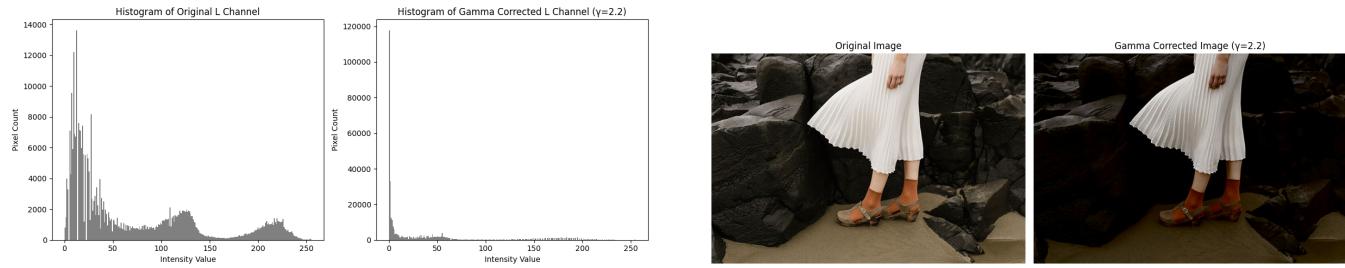
color information, which makes it very useful for image processing tasks.

Why Apply Gamma Correction Only on the L Channel? The L channel contains the lightness (brightn information separated from color so it makes sense to apply brightness adjustments there.

```

1 # Convert BGR to Lab
2 lab_img = cv2.cvtColor(img3, cv2.COLOR_BGR2LAB)
3 # Split the Lab image into L, a, b channels
4 L, a, b = cv2.split(lab_img)
5 #Normalize only L channel to [0,1] for gamma correction
6 L_norm = L / 255.0
7 # choose the gamma value = 2.2
8 gamma = 2.2
9 L_gamma = np.power(L_norm, gamma)
10 L_corrected = np.uint8(L_gamma * 255)
11 lab_corrected = cv2.merge((L_corrected, a, b)) #merge
12 result = cv2.cvtColor(lab_corrected, cv2.COLOR_LAB2BGR)

```



4. Increasing the vibrance of a photograph is probably achieved by applying an intensity transformation such as

$$f(x) = \min \left(x + a \times 128 e^{-\frac{(x-128)^2}{2\sigma^2}}, 255 \right)$$

to the saturation plane, where x is the input intensity, a [0,1] and $\sigma = 70$.

(a) Split the image into hue, saturation, and value planes.

```

1 # Convert the image from BGR to HSV color space
2 hsv_img = cv2.cvtColor(img4, cv2.COLOR_BGR2HSV)
3 # Split the HSV image into individual channels
4 h, s, v = cv2.split(hsv_img)

```

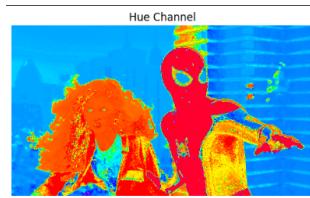
(b) Apply the aforementioned intensity transformation to the saturation plane.

```

1 # Parameters
2 a = [0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1]
3 sigma = 70.0
4
5 for i in a:
6     # Calculate the Gaussian term
7     gaussian_term = i * 128 * np.exp(-((s.astype(np.float32) - 128) ** 2) / (2 * (sigma
8         ** 2)))
9     # Apply the function f(x) = min(x + gaussian_term, 255)
10    s_transformed = s.astype(np.float32) + gaussian_term
11    s_transformed = np.clip(s_transformed, 0, 255).astype(np.uint8)
12    # Recombine three channels
13    hsv_transformed = cv2.merge([h, s_transformed, v])
14    # Convert back to BGR to display
15    result = cv2.cvtColor(hsv_transformed, cv2.COLOR_HSV2BGR)

```

Saturation Plane Transformed



Saturation Channel



Value Channel



5. Write a function of your own to carry out histogram equalization on the image shown in Fig. 5. Show the histograms before and after equalization.

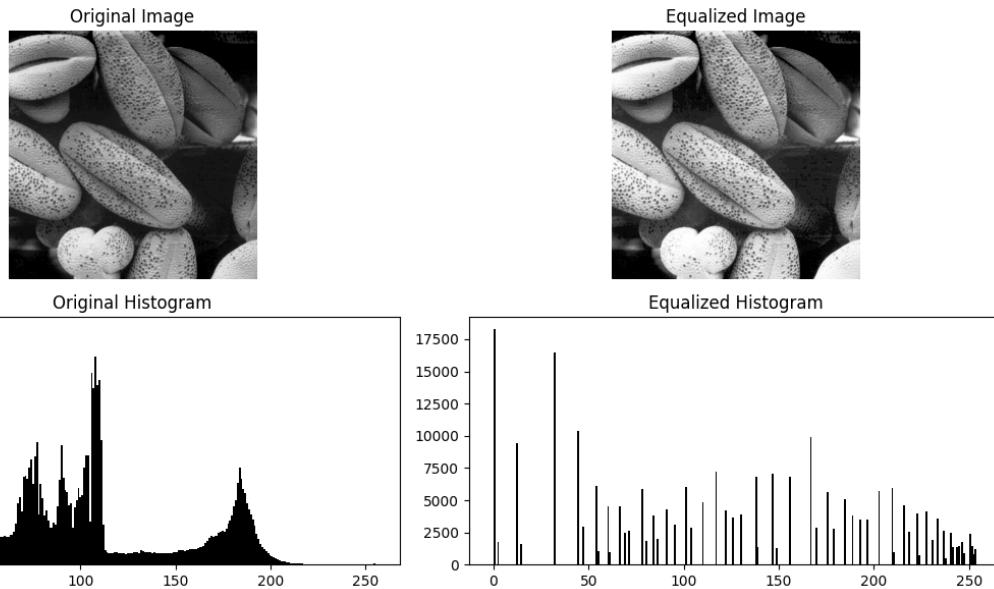
```

1 def histogram_equalization(image):
2     # Flatten image array and calculate histogram
3     hist, bins = np.histogram(image.flatten(), 256, [0,256])
4     # Calculate the cumulative distribution function (CDF)
5     cdf = hist.cumsum()
6     # Mask all zero values (to avoid division by zero)
7     cdf_masked = np.ma.masked_equal(cdf, 0)
8     # Normalize the CDF
9     cdf_min = cdf_masked.min()
10    cdf_max = cdf_masked.max()
11    cdf_norm = (cdf_masked - cdf_min) * 255 / (cdf_max - cdf_min)
12    # Fill masked values with zero
13    cdf_final = np.ma.filled(cdf_norm, 0).astype('uint8')
14    # Map the original image pixels to equalized values using the normalized CDF as a
15        lookup table
16    img_equalized = cdf_final[image]
17    return img_equalized

```

6. In this question, we will apply histogram equalization only to the foreground of an image to produce an image with a histogram equalized foreground

(a) Open the image in Fig. 6, split it into hue, saturation, and values and display these planes in grayscale



```

1 img_rgb = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2RGB)
2 hsv = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2HSV)
3 h, s, v = cv2.split(hsv)

```

(b) Select the appropriate plane to threshold in extract the foreground mask. A mask is a binary image.
(c) Now obtain the foreground only using `cv.bitwise_and` and compute the histogram

```

1 # saturation threshold (keep moderately high saturation as foreground)
2 s_lower, s_upper = 60, 255
3 mask_s = cv2.inRange(s, s_lower, s_upper) # binary mask from S
4
5 # Optionally refine using morphology to clean small noise and fill small holes
6 kernel = np.ones((3,3), np.uint8)
7 mask_bin = cv2.morphologyEx(mask_s, cv2.MORPH_OPEN, kernel, iterations=1)
8 mask_bin = cv2.morphologyEx(mask_bin, cv2.MORPH_CLOSE, kernel, iterations=1)

```

```

1 # Extract foreground from RGB (visualization) and V (for intensity histogram)
2 foreground_rgb = cv2.bitwise_and(img_rgb, img_rgb, mask=mask_bin)
3
4 # Choose intensity channel to analyze for histogram equalization:
5 # Typically use the Value (V) channel in HSV for brightness equalization.
6 foreground_v = cv2.bitwise_and(v, v, mask=mask_bin)
7
8 # Compute histogram only over foreground pixels
9 # Note: foreground_v already zeros out background; we can hist over the non-zero mask
10 hist_fg, bins = np.histogram(foreground_v[mask01==1], bins=256, range=(0,256))

```

(d) Obtain the cumulative sum of the histogram using `np.cumsum`.

```

1 cdf_fg = np.cumsum(hist_fg) # length 256
2 # Normalize CDF to [0, 1]
3 cdf_min = cdf_fg[np.nonzero(cdf_fg)].min() if np.any(cdf_fg) else 0
4 N_fg = cdf_fg[-1] if len(cdf_fg) > 0 else 1 # total foreground pixels

```

(e) Use the formulas in slides to histogram-equalize the foreground

```

1 L = 256
2 # Build mapping for 0..255
3 equal_map = np.zeros(256, dtype=np.uint8)
4 if N_fg > 0 and cdf_min < N_fg:
5     equal_map = np.round((cdf_fg - cdf_min) / (N_fg - cdf_min) * (L - 1)).clip(0, 255).astype(np.uint8)
6
7 # Apply mapping only to foreground pixels on the V channel
8 v_eq = v.copy()
9 # Map: for foreground pixels, replace v by equalized value, else keep original v
10 # To avoid touching background, we only remap at mask==1
11 v_eq[mask01==1] = equal_map[v[mask01==1]]
12
13 # Histogram after equalization for foreground pixels
14 hist_fg_eq, _ = np.histogram(v_eq[mask01==1], bins=256, range=(0,256))

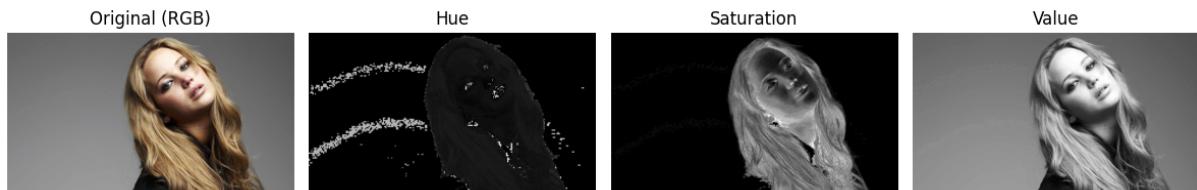
```

(f) Extract the background and add with the histogram equalized foreground.

```

1 hsv_eq = hsv.copy()
2 hsv_eq[:, :, 2] = v_eq
3
4 # Convert back to RGB for visualization
5 img_eq_rgb = cv2.cvtColor(hsv_eq, cv2.COLOR_HSV2RGB)
6
7 # Alternatively, explicitly recompose using the mask:
8 # - Foreground from img_eq_rgb (equalized V)
9 # - Background from original img_rgb (unchanged)
10 foreground_eq_rgb = (img_eq_rgb * mask01[:, None]).astype(np.uint8)
11 background_rgb = (img_rgb * (1 - mask01[:, None])).astype(np.uint8)
12 recomposed_rgb = (foreground_eq_rgb + background_rgb).astype(np.uint8)

```



7. Filtering with the Sobel operator can compute the gradient.

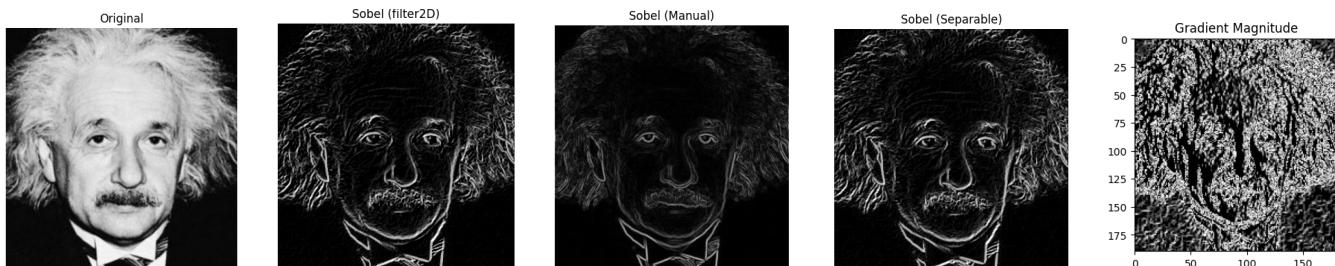
(a) Using the existing filter2D to Sobel filter the image.

(b) Write your own code to Sobel filter the image

```

1 # Sobel X separable filters
2 #Sobel filter using cv2.filter2D
3 sobel_x_kernel = np.array([[1, 0, -1],
4                             [2, 0, -2],
5                             [1, 0, -1]], dtype=np.float32)
6 sobel_y_kernel = np.array([[1, 2, 1],
7                             [0, 0, 0],
8                             [-1, -2, -1]], dtype=np.float32)
9 sobel_x_a = cv2.filter2D(img7, -1, sobel_x_kernel)
10 sobel_y_a = cv2.filter2D(img7, -1, sobel_y_kernel)
11 sobel_a = cv2.magnitude(sobel_x_a.astype(np.float32), sobel_y_a.astype(np.float32))
12 # My Sobel filter
13 def my_convolution(image, kernel):
14     kh, kw = kernel.shape
15     pad_h, pad_w = kh//2, kw//2
16     padded = np.pad(image, ((pad_h, pad_h), (pad_w, pad_w)), mode='constant')
17     output = np.zeros_like(image, dtype=np.float32)
18
19     for i in range(image.shape[0]):
20         for j in range(image.shape[1]):
21             region = padded[i:i+kh, j:j+kw]
22             output[i, j] = np.sum(region * kernel)
23     return output
24 sobel_x_b = my_convolution(img7, sobel_x_kernel)
25 sobel_y_b = my_convolution(img7, sobel_y_kernel)
26 sobel_b = np.sqrt(sobel_x_b**2 + sobel_y_b**2)
27 # Separable convolution
28 gx_row = np.array([[1, 0, -1]], dtype=np.float32) # horizontal derivative
29 gx_col = np.array([[1], [2], [1]], dtype=np.float32) # smoothing vertically
30 # X-direction
31 temp_x = cv2.filter2D(img7, -1, gx_row)
32 sobel_x_c = cv2.filter2D(temp_x, -1, gx_col)
33 # Y-direction (transpose order)
34 gy_row = np.array([[1, 2, 1]], dtype=np.float32)
35 gy_col = np.array([[1], [0], [-1]], dtype=np.float32)
36 temp_y = cv2.filter2D(img7, -1, gy_col)
37 sobel_y_c = cv2.filter2D(temp_y, -1, gy_row)
38 sobel_c = cv2.magnitude(sobel_x_c.astype(np.float32), sobel_y_c.astype(np.float32))

```



8. Write a program to zoom images by a given factor s ($0,10]$. You must use a function to zoom the image, which can handle (a) nearest-neighbor, and (b) bilinear interpolation.

```

1 | def zoom_image(img, scale_factor, method='nearest'):

```

```

2     height, width = img.shape[:2]
3     new_height, new_width = int(height * scale_factor), int(width * scale_factor)
4     if method == 'nearest':
5         interp = cv2.INTER_NEAREST
6     elif method == 'bilinear':
7         interp = cv2.INTER_LINEAR
8     else:
9         raise ValueError("Method must be 'nearest' or 'bilinear'")
10    return cv2.resize(img, (new_width, new_height), interpolation=interp)

```

```

1 def normalized_ssd(img1, img2):
2     # Ensure both are float for computation
3     img1 = img1.astype(np.float32)
4     img2 = img2.astype(np.float32)
5     # If different sizes, compute over overlapping region
6     h = min(img1.shape[0], img2.shape[0])
7     w = min(img1.shape[1], img2.shape[1])
8     img1_c = img1[:h, :w]
9     img2_c = img2[:h, :w]
10    diff = img1_c - img2_c
11    ssd = np.sum(diff * diff)
12    denom = np.sum(img1_c * img1_c) + 1e-12 # avoid divide-by-zero
13    return ssd / denom

```

For image 1 - Normalized SSD (nearest neighbor): 0.012057948510772873 and Normalized SSD (bilinear): 0.010184055645914382



9. (a) Use grabCut to segment the image. Show the final segmentation mask, foreground image, and background image.

```

1 # rectangle init
2 rect = tuple(map(int, (0.1*w_, 0.1*h_, 0.8*w_, 0.8*h_)))
3 # grabcut (rect)
4 mask = np.zeros(img_rgb.shape[:2], np.uint8)
5 bldModel = np.zeros((1,65), np.float64)
6 fgdModel = np.zeros((1,65), np.float64)
7 cv2.grabCut(img_bgr, mask, rect, bldModel, fgdModel, 5, cv2.GC_INIT_WITH_RECT)
8 # yellow refinement
9 hsv = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2HSV)
10 yellow = cv2.inRange(hsv, (15,60,80), (35,255,255))
11 mask_refined = mask.copy()
12 sel = yellow > 0
13 mask_refined[sel & (mask_refined == 0)] = cv2.GC_PR_FGD # 3
14 # grabcut (mask)
15 cv2.grabCut(img_bgr, mask_refined, None, bldModel, fgdModel, 5, cv2.GC_INIT_WITH_MASK)
16 # final mask (0/1)
17 final_mask = np.where((mask_refined == cv2.GC_FGD) | (mask_refined == cv2.GC_PR_FGD), 1,
0).astype(np.uint8)

```

```

18 # compose fg/bg
19 foreground = (img_rgb * final_mask[..., None]).astype(np.uint8)
20 background = (img_rgb * (1 - final_mask[..., None])).astype(np.uint8)

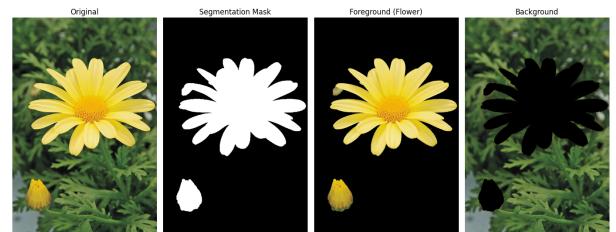
```

(b) Produce an enhanced image with a substantially blurred background. Display the original image alongside the enhanced image.

```

1 def feather_mask(binary_mask, feather_px=5):
2     # Feather/soften edges by Gaussian blur on a float mask
3     m = binary_mask.astype(np.float32)
4     m = cv2.GaussianBlur(m, (2*feather_px+1, 2*feather_px+1), feather_px)
5     m = np.clip(m, 0, 1)
6     return m
7
8 # Create a strongly blurred version
9 blur_ksize = 41 # strong blur; must be odd
10 blur_sigma = 0
11 blurred_bgr = cv2.GaussianBlur(img_bgr, (blur_ksize, blur_ksize), blur_sigma)
12 blurred_rgb = cv2.cvtColor(blurred_bgr, cv2.COLOR_BGR2RGB)
13
14 # Feather the mask for smoother compositing
15 soft_mask = feather_mask(final_mask, feather_px=5)[..., None] # HxWx1
16
17 # Composite: keep foreground sharp, background blurred
18 enhanced_rgb = (img_rgb.astype(np.float32) * soft_mask + blurred_rgb.astype(np.float32) *
    (1 - soft_mask)).astype(np.uint8)

```



(c) Why is the background just beyond the edge of the flower quite dark in the enhanced image?

- Mixed pixels at boundaries: Real-world edges contain pixels that are a mixture of foreground and background due to sensor sampling and lens point spread; when multiplied by a binary/near-binary mask, these partially mixed pixels can lose brightness and appear darker.

github repo